




Container Basics

what why secure?

n3 | {π0s!s

- day time cybersecurity engineer
- night time student
 - perpetual newb
- tweet @n3krosis11
- email n3krosis@protonmail.com

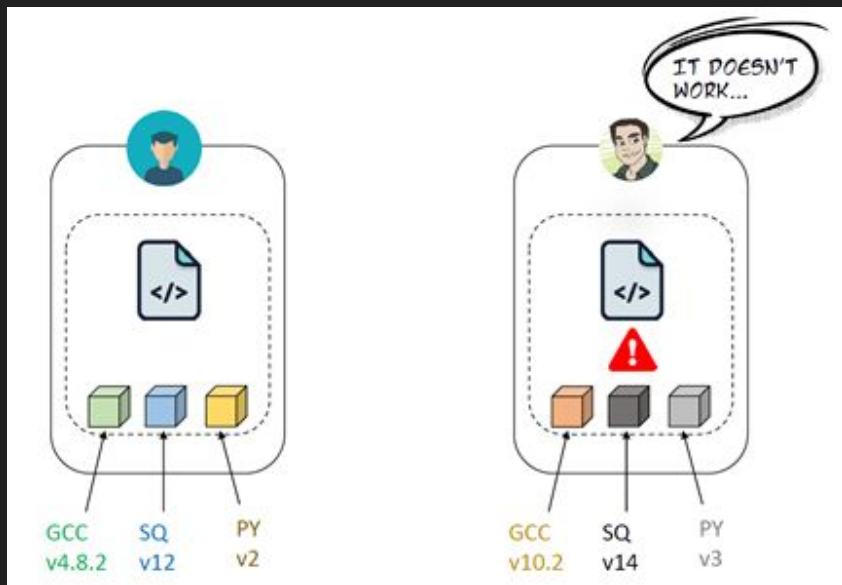


today's lecture

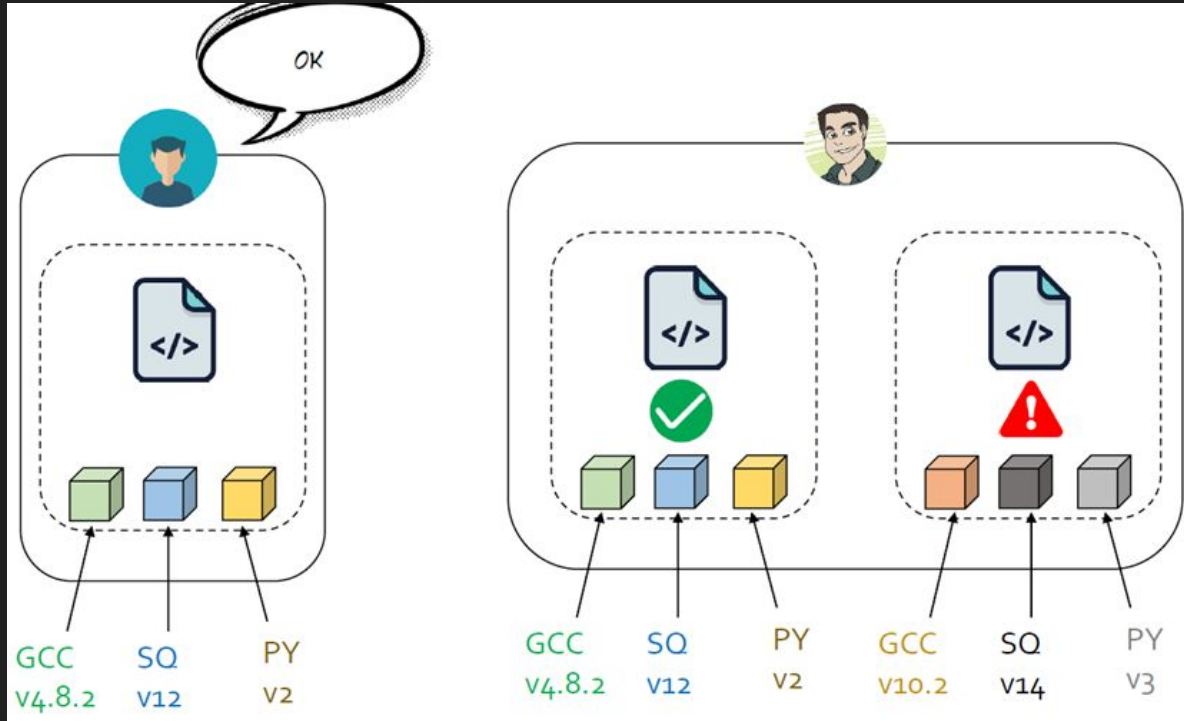
- content adapted from class lecture and blackhat talk
 - Shelby Thomas Ph.D. @realshelbyt on Twitter
 - Brandon Edwards & Nick Freeman
- what are container
- container platforms
- docker demo
- security
- escape demo
- questions

what can containers do?

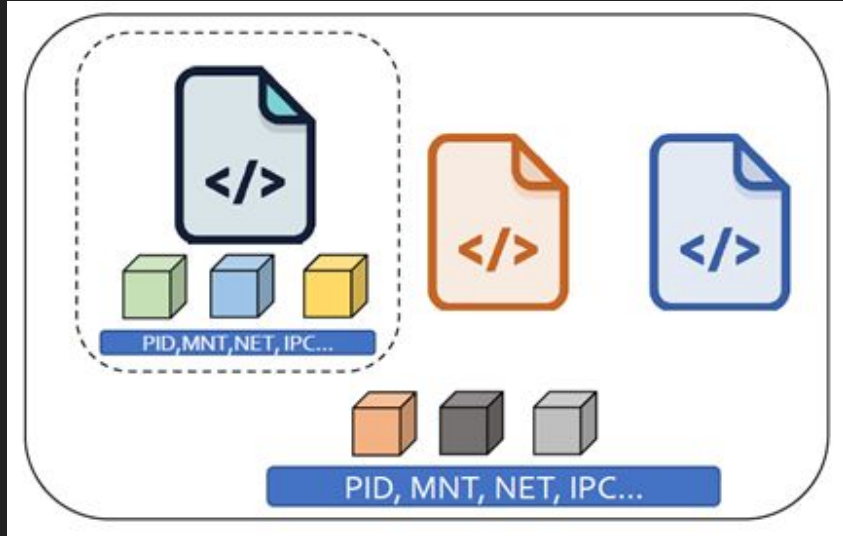
- send code to another
- different development environments
- prevents everything breaking
- can't easily send computers



multiple versions & dependencies

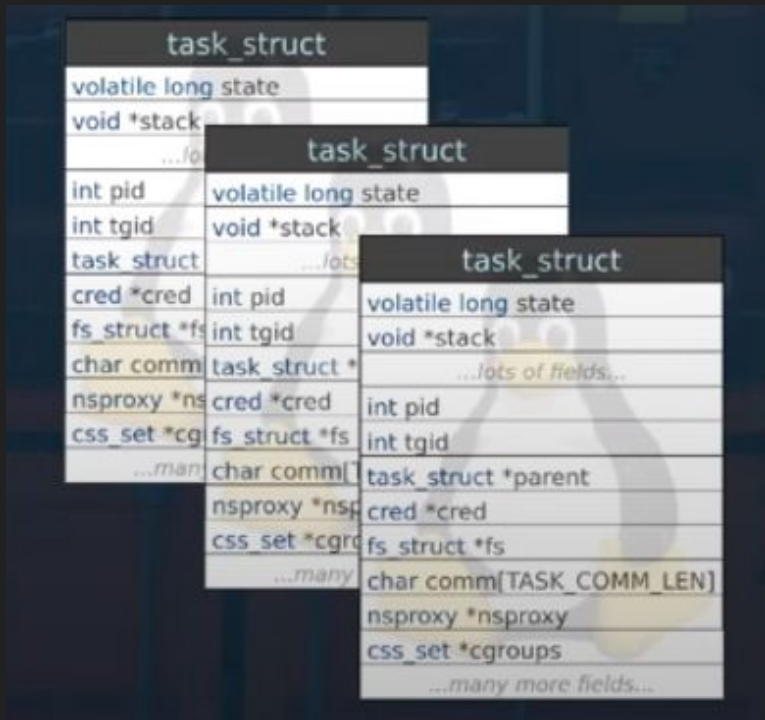


what is a container?



- containers are a type of virtualization for linux kernel
- what kernel capabilities might a process (program) need?
- the capabilities are called namespaces
- new containers mean new namespaces


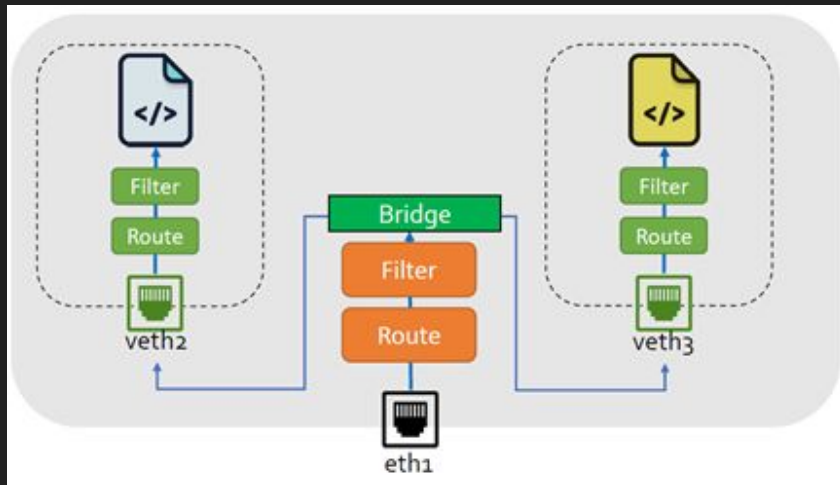
little deeper, what is a container?



- task struct with key features
 - credentials
 - user permissions
 - capabilities
 - separate root things
 - filesystem
 - container's own fs
 - namespaces
 - define things it can do
 - cgroups
 - resource limitations
 - linux security modules
 - apparmor, selinux
 - seccomp
 - restrict syscalls

closer look at the network namespace

- what is a network?
- network namespace isolation
- operational benefit: experiment with rules because they are tied to a container
- often managed through a virtual bridge/switch



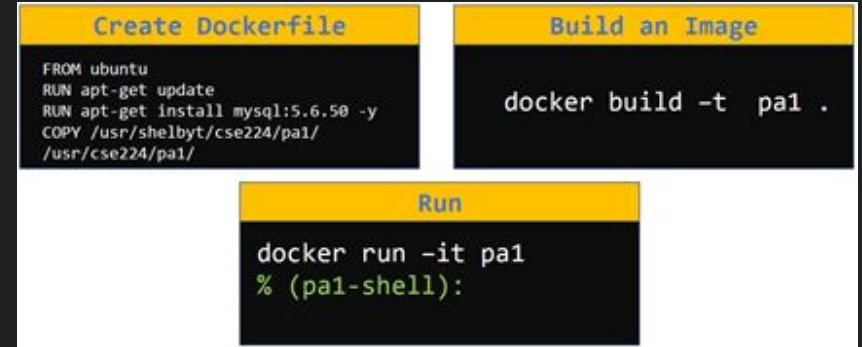
container platforms

```
$ ip link add veth0 type veth peer name veth1
$ ip link set veth1 netns pa1
$ ip netns exec pa1 ip addr add 10.1.1.1/24 dev veth1
$ ip netns exec pa1 ip link set dev veth1 up
```

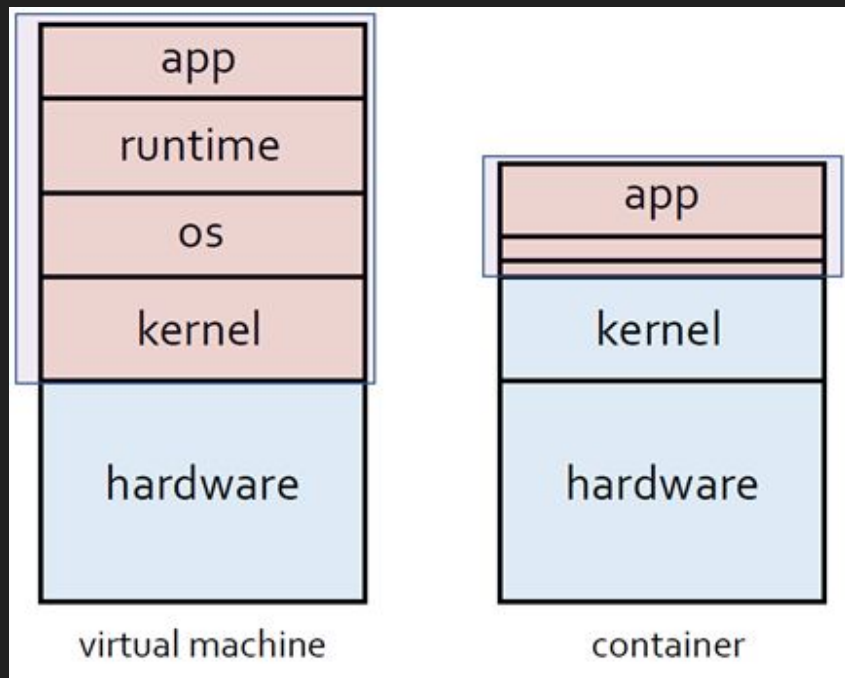
- these are all kernel capabilities (i.e., can build a container without additional software)
- struggles of creating namespaces
- repeat for additional isolation



- container software makes it easy to manage containers and namespaces
- docker abstracts underlying namespaces
- isolated shell, no manual configuration of namespaces needed



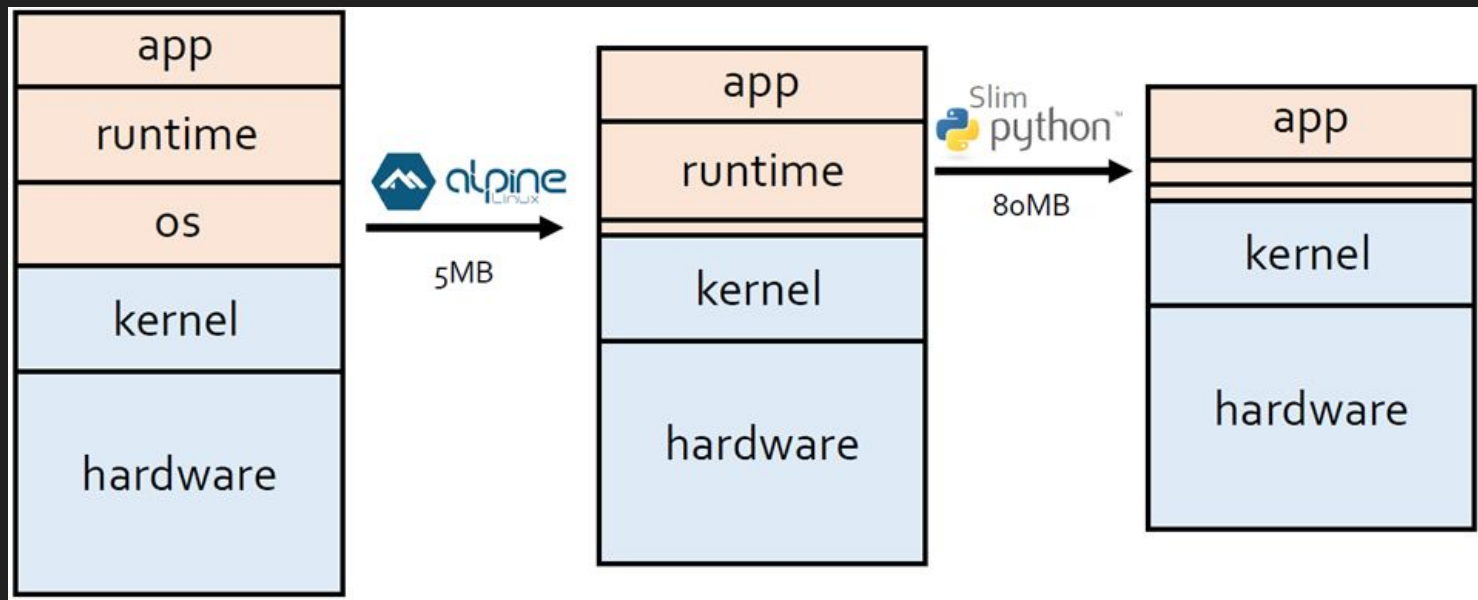
virtual machine vs. containers



- different isolation points
- container benefits
 - lightweight (megabytes vs gigabytes [10x])
 - fast startup (milliseconds vs. minutes)
 - simple to build, deploy, send, & maintain


virtual machine vs. containers

- benefits come from specialization




serverless computing

- os and runtime are partially managed by cloud providers
- serverless
 - on demand
 - focus on application
 - subject to fluctuations
 - questionable security
- traditional cloud
 - persistent
 - more control
 - fast and predictable
 - difficult to manage



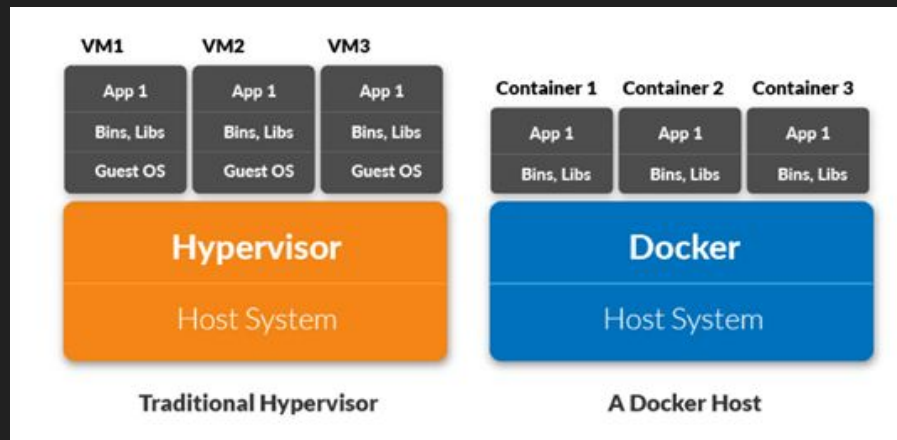
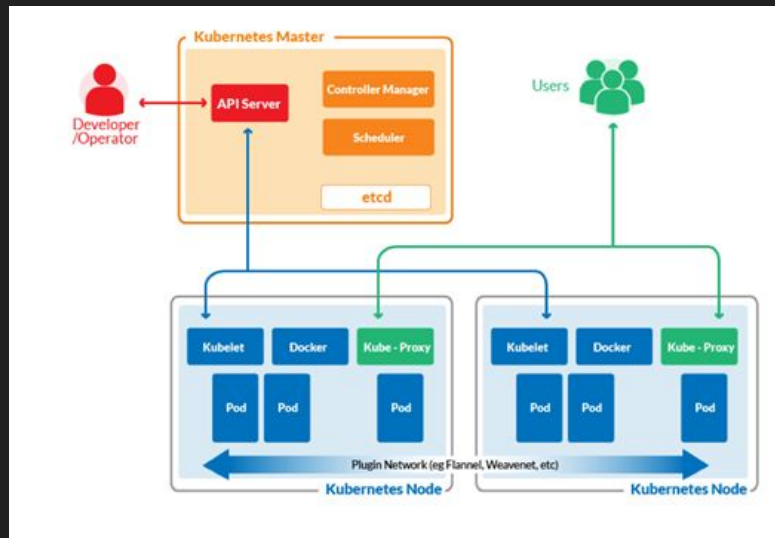
kubernetes aside



- Automates operational tasks of container management
 - Automated operations
 - Infrastructure abstraction
 - Service health monitoring
- vs. docker
 - kubernetes == architecture
 - docker == host application

kubernetes vs. docker

kubernetes

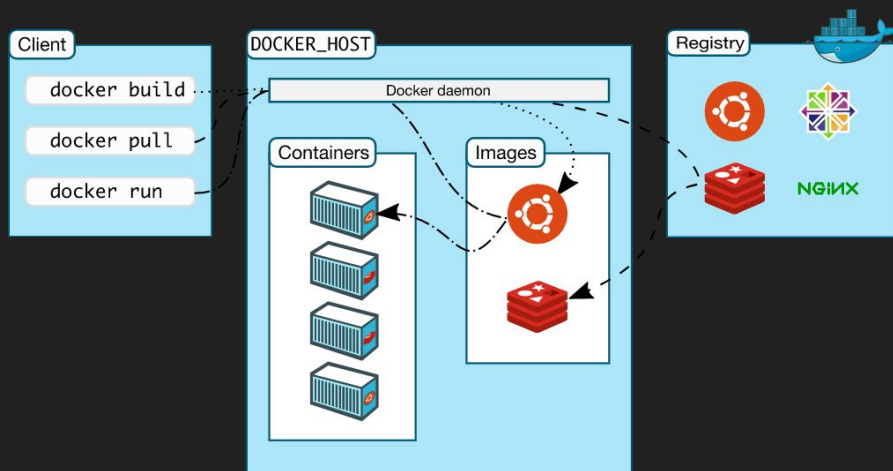


docker

setup demo



docker basic operation



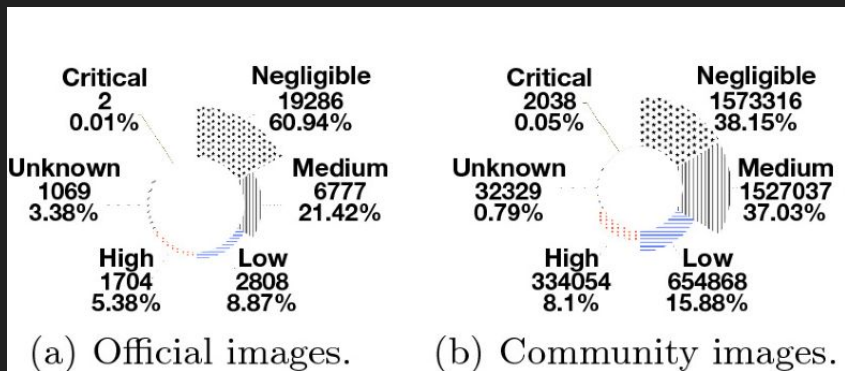
- **format**
 - `docker cmd [IMAGE_NAME]`
 - `docker cmd [CONTAINER_NAME]`
- **popular**
 - `create / run [image]`
 - `start / stop [container]`
 - `restart [container]`
 - `pause / unpause [container]`
 - `ps / ps -a`
 - `rm`

dockerfile + compose

- dockerfile
 - text document that contains all the commands a user could call on the command line to assemble an image
 - **docker build** allows for automated build that executes several command-line instructions in succession
- compose
 - tool for defining and running multi-container docker applications
 - e.g. development environments, automated testing environments, single host deployments

```
# syntax=docker/dockerfile:1
FROM python:3.7-alpine
WORKDIR /code
ENV FLASK_APP=app.py
ENV FLASK_RUN_HOST=0.0.0.0
RUN apk add --no-cache gcc musl-dev linux-headers
COPY requirements.txt requirements.txt
RUN pip install -r requirements.txt
EXPOSE 5000
COPY . .
CMD ["flask", "run"]
```


```
version: "3.9"
services:
  web:
    build: .
    ports:
      - "5000:5000"
  redis:
    image: "redis:alpine"
```



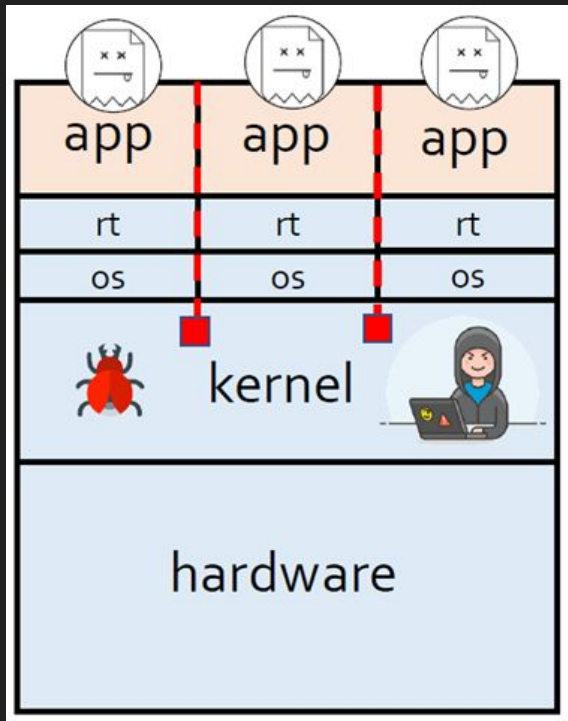
from Understanding the Security
Risks of Docker Hub 09-2020

- benefits
 - image repos
 - teams and orgs
 - github bitbucket integration
 - automated builds
 - webhooks
 - official and publisher images
- security issues
 - trust everything you download?
 - images contain patched vulns
 - malicious images
 - ffmpeg -> btc mining
 - Dec 2020 - reported 51% of the Docker Hub images had exploitable vulnerabilities
 - company site offline

but is it safe

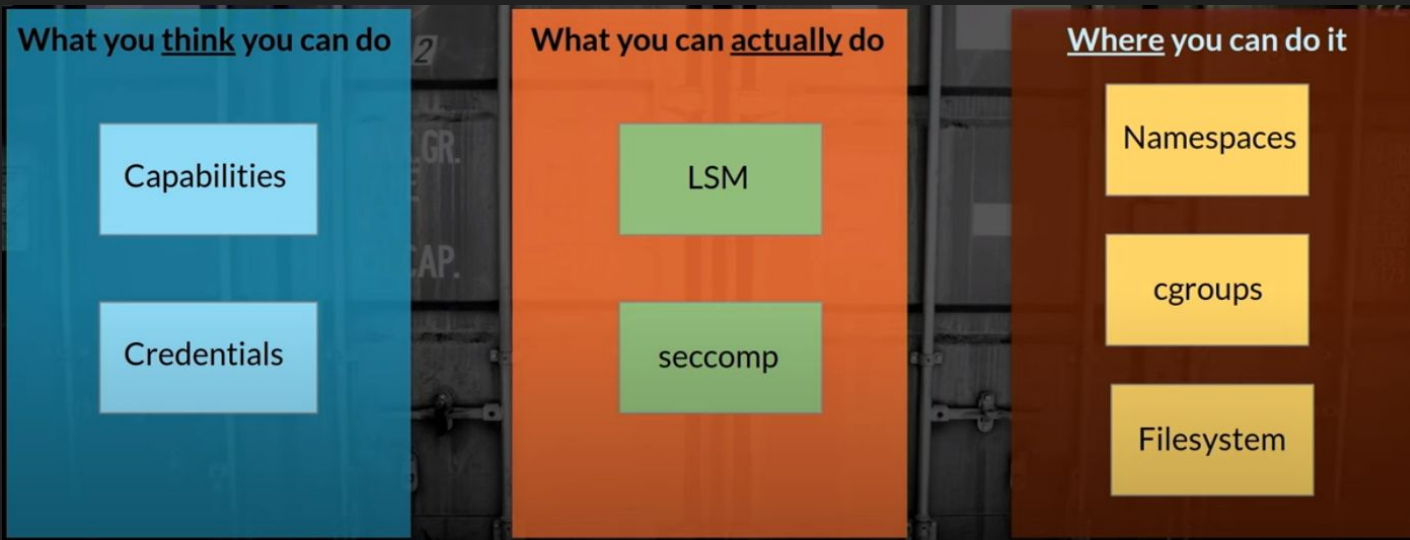


security layers



- container image itself and the software inside
- interaction between container, host os, and other containers on the same host
- host os
- networking and storage
- security at runtime

container security model



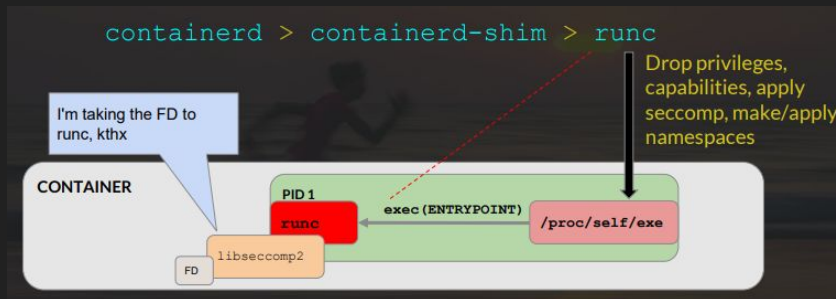
recent cves

- CVE-2019-5736
 - CVSS 3.x -- 8.6 HIGH
 - CVSS:3.1/AV:L/AC:L/PR:N/UI:R/S:C/C:H/I:H/A:H
 - runc 1.0-rc6, Docker before 18.09.2
 - attackers overwrite host runc binary and consequently obtain host root access
- 9 CVEs since <= 7.5 Medium

regular



escape



brandon edwards & nick freeman

**Engine bugs are awesome,
but probably not how you'll get popped**

common weaknesses



- exposed docker sockets
 - docker cmds, use curl too
 - inception (docker in docker)
- privileged containers
 - init with `--privileged` flag
- excessive capabilities
 - privileged or root
- sensitive mounts
 - access to hosts fs
- kernel exploits
 - dirty cow
 - escape to update struct
 - credentials
 - namespaces
 - fs struct

escape patterns


- step 1:
 - memory layout, state grooming, etc.
- step 2:
 - trigger bug for kernel access
- step 3:
 - rop to disable smep/smap
- step 4:
 - return to userland
- step 5:
 - `commit_creds(\prepare_kernel_creds(0));`



escape demo



revised container security model



securing containers -- owasp

- RULE #0 - Keep Host and Docker up to date
- RULE #1 - Do not expose the Docker daemon socket (even to the containers)
- RULE #2 - Set a user
- RULE #3 - Limit capabilities
- RULE #4 - Add
-no-new-privileges flag
- RULE #5 - Disable
inter-container communication
- RULE #6 - Use Linux Security Module
- RULE #7 - Limit resources
- RULE #8 - Set filesystem and volumes to read-only
- RULE #9 - Use static analysis tools
- RULE #10 - Set the logging level to at least INFO
- Rule #11 - Lint the Dockerfile at build time

automate security

- Docker Bench for Security
 - script that checks for dozens of common best-practices around deploying Docker containers in production
 - tests are all automated
 - inspired by the CIS Docker Benchmark v1.2.0
- Ansible from DoD
 - define and enforce system and application configurations
 - configurations that implement most of the Docker Enterprise 2.x Linux/Unix STIG



docker bench demo



questions



further reading

- Kubernetes vs. Docker: A Primer
 - <https://containerjournal.com/topics/container-systems/kubernetes-vs-docker-a-primer/>
- Installing Docker on Debian 10
 - <https://www.digitalocean.com/community/tutorials/how-to-install-and-use-docker-on-debian-10>
- Container Security For Development Teams
 - <https://snyk.io/learn/container-security>
- A Compendium of Container Escapes (Black Hat 2019)
 - <https://youtu.be/BQlqita2D2s>
- Docker Security Cheat Sheet
 - https://cheatsheetseries.owasp.org/cheatsheets/Docker_Security_Cheat_Sheet.html

further reading

- Understanding Docker container escapes
 - <https://blog.trailofbits.com/2019/07/19/understanding-docker-container-escapes/>
- Vulnerable Docker containers for testing
 - <https://hub.docker.com/u/vulnerables>
- Example of Docker escape
 - <https://github.com/Swordfish-Security/Pentest-In-Docker>
- Understanding the Security Risks of Docker Hub
 - https://link.springer.com/chapter/10.1007/978-3-030-58951-6_13